

MSI/Git Workshop

Tom Kono

2017-11-14

Slides: <https://z.umn.edu/msi-git>

© 2017 Regents of the University of Minnesota. All rights reserved.



Outline

- Part 1: Intro to MSI Systems
 - HPC systems, Storage, Queues, Modules
- Part 2: Parallelization and Task Arrays
 - GNU Parallel, PBS Task Arrays
- Part 3: Git and GitHub

© 2017 Regents of the University of Minnesota. All rights reserved.



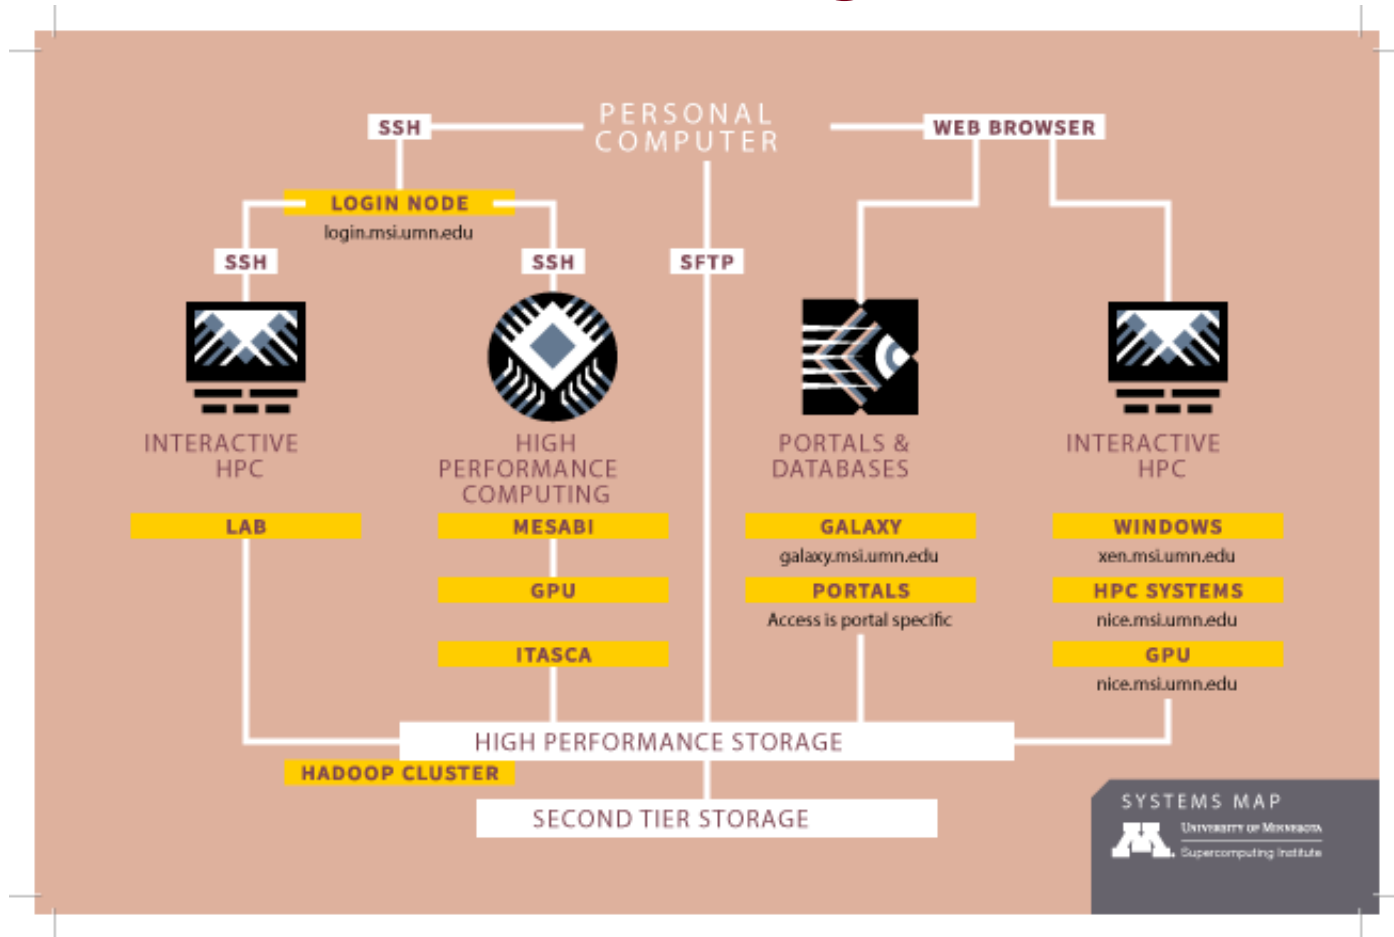
Part 1: MSI Organization

- Divisions of computing support:
 - **High performance computing (HPC)**
 - **Interactive HPC (Linux)**
 - **Storage systems**
 - GPU nodes

© 2017 Regents of the University of Minnesota. All rights reserved.



Part 1: MSI Organization



© 2017 Regents of the University of Minnesota. All rights reserved.

Part 1: HPC Overview

- Non-interactive jobs: all work must be done as shell scripts
- Run CentOS Linux
- Hundreds of cores, hundreds of GB RAM, days to weeks of runtime
- Two clusters: Itasca and Mesabi

© 2017 Regents of the University of Minnesota. All rights reserved.



Part 1: HPC - Itasca

- Many 22 GB RAM, 8 core nodes, no node sharing
- Several 64+ GB RAM, 16 core nodes
- Great for short jobs that can be run in parallel

© 2017 Regents of the University of Minnesota. All rights reserved.



Part 1: HPC - Mesabi

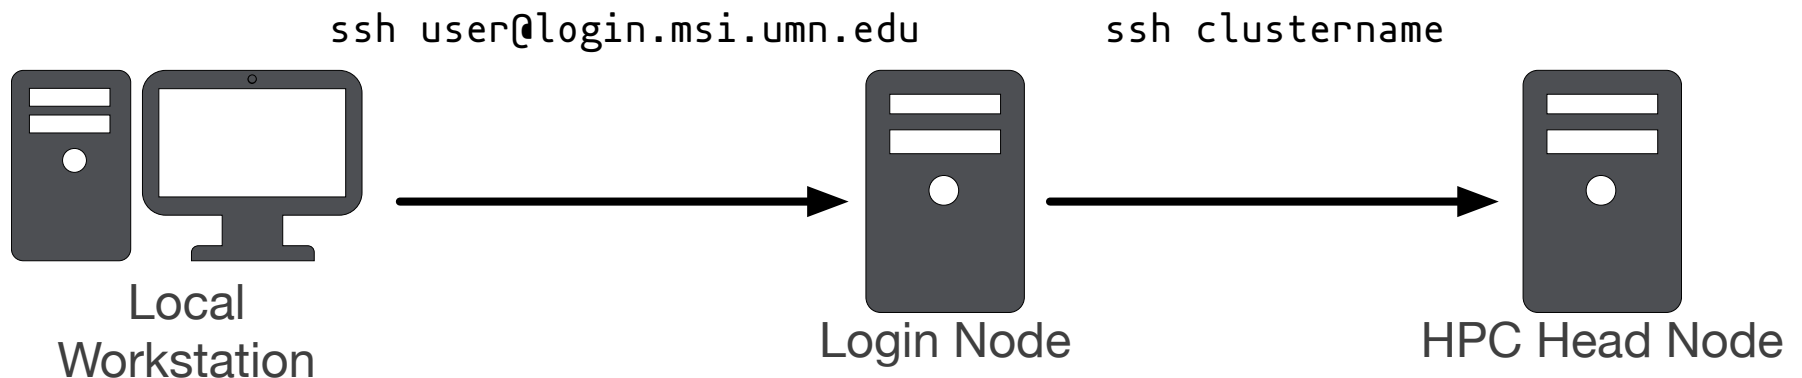
- Very diverse system: long-running jobs, highly parallel jobs, high-memory jobs
- Has some GPU processing capacity

© 2017 Regents of the University of Minnesota. All rights reserved.



Part 1: Accessing HPC Systems

- Must be done on-campus, or through university VPN
- Access via terminal and `ssh`:



© 2017 Regents of the University of Minnesota. All rights reserved.

Part 1: HPC Queues

- Scheduling jobs is done by TORQUE
- Jobs must be submitted from the cluster head node
- You can submit to a variety of job queues, depending on the parameters of the analysis

© 2017 Regents of the University of Minnesota. All rights reserved.



Part 1: HPC Queues

- Itasca queues:

Queue name	Number of Nodes	Processor cores per node	Wallclock Limit	Total Node Memory Limit	Per-core Memory Limit	Local Scratch (GB/node)	Per User Running Jobs	Per User Idle Jobs (gaining priority in queue)
batch (default)	1086 nodes (8688 cores)	8	24 hours	22gb	2750mb	90 GB	60	8
devel	32 nodes (256 cores)	8	2 hours	22gb	2750mb	90 GB		
long	28 nodes (224 cores)	8	48 hours	22gb	2750mb	90 GB		
sb	35 nodes (560 cores)	16	48 hours	62gb	3875mb	112 GB		
sb128	8 nodes (128 cores)	16	96 hours	126gb	7875mb	534 GB		
sb256	8 nodes (128 cores)	16	96 hours	254gb	15875mb	534 GB		

Service Unit (SU) rate: 1.5 CPU hours / SU

© 2017 Regents of the University of Minnesota. All rights reserved.

Part 1: HPC Queues

- Mesabi queues:

Queue name	Node Sharing	Max Nodes Per Job	Min Nodes Per Job	Processor cores per node	Wallclock Limit	Total Node Memory Limit	Per-core Memory Limit	Local Scratch (GB/node)	Per User Limits	Per Group Limits
small [€]	Yes	9	None	24	96 hours	62gb	2580mb	390 GB	500 Jobs	1800 total cores [‡]
large	No	48	10	24	24 hours	62gb	2580mb	390 GB	4 Jobs	16 Jobs
widest [§]	No	360	49	24	24 hours	62gb	2580mb	390 GB	4 Jobs	16 Jobs
max	Yes	1 (single core per job)	None	1	696 hours	62gb	62gb	390 GB	4 Jobs	16 Jobs
ram256g	Yes	2	None	24	96 hours	252gb	10580mb	390 GB	2 nodes	1800 total cores [‡]
ram1t	Yes	2	None	32 [†]	96 hours	998gb	31180mb	228 GB	2 nodes	1800 total cores [‡]
k40 GPU nodes [*]	No	40	None	24	24 hours	126gb	5290mb	390 GB	None	1800 total cores [‡]
mesabi (default)	The mesabi queue is a meta-queue, which will automatically route jobs to the small , large , widest , or max queues, according to where each job will best fit based on the resource request.									

Service Unit (SU) rate: 1.5 CPU hours / SU

© 2017 Regents of the University of Minnesota. All rights reserved.

Part 1: Interactive HPC

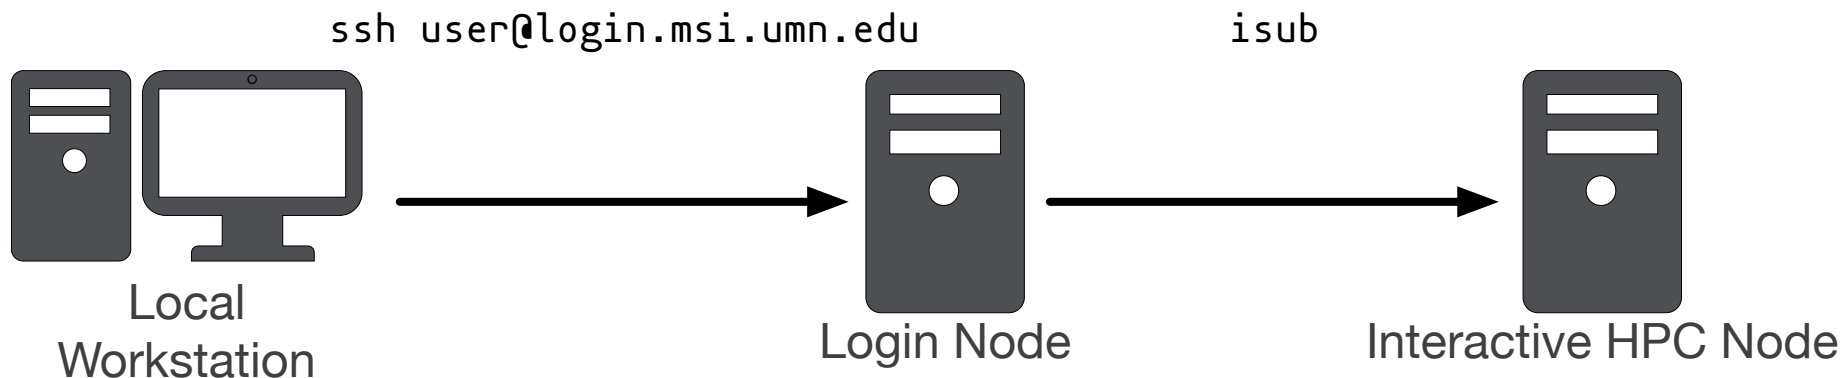
- Small interactive jobs: some GUI applications, R, MATLAB
- 1-16 cores, tens of GB RAM, hours of runtime
- Useful for testing scripts and programs before throwing them onto HPC

© 2017 Regents of the University of Minnesota. All rights reserved.



Part 1: Accessing Interactive HPC

- Must be done on-campus, or through university VPN
- Terminal, `ssh`, and `isub`:



© 2017 Regents of the University of Minnesota. All rights reserved.

Part 1: Interactive HPC

- Use interactive HPC if you want to run or test analysis software
 - You can technically run software on the HPC head node, but your processes can (and will) be killed
- Never run analyses on the login node. You can lock up access to the systems and people will yell at you.

© 2017 Regents of the University of Minnesota. All rights reserved.



Part 1: Storage Systems

- **Primary:** Home directories on SSD. Snapshot backups, quotas enforced
- **Scratch:** Temp storage, quotas not enforced, files kept for at most 30 days
- **Tier 2:** Long-term storage, Amazon S3 system, no snapshots.
- **Archive:** Tape storage, system in beta

© 2017 Regents of the University of Minnesota. All rights reserved.



Part 1: Important Directories

- Primary storage

`/home/groupname/username`: Home directory. Scripts/small data files

`/home/groupname/shared`: Group shared directory. Software, reference sequences, common datasets.

- Scratch storage

`/panfs/roc/scratch`: Global scratch storage.

`/scratch.local`: Node local scratch storage. Cleaned out at end of job.

© 2017 Regents of the University of Minnesota. All rights reserved.

Part 1: Storage Guidelines

- Primary has a group quota. Once quota is met, **no new files can be made.**
- Remove files from scratch when done.
- **Tier 2 is unlimited, but not backed up. Be careful.**
- Archive is limited access, you'll know if you have it

© 2017 Regents of the University of Minnesota. All rights reserved.



Part 1: Modules

- Software on MSI is loaded using the `module` system
- Many common pieces of bioinformatics software (bowtie2, bwa, samtools, blast, etc.) are available as modules
- `module avail` to list, `module load` to use

© 2017 Regents of the University of Minnesota. All rights reserved.

Part 1: Modules

- You can download/compile software if there is no module available for it.
 - Mostly unsupported: you should know what you are doing
 - You can request software to be installed on the system; email the help desk.
- **Always specify a version when loading a module, for reproducibility.**

© 2017 Regents of the University of Minnesota. All rights reserved.

Part 1: Let's Access MSI

1. Open terminal program
 - Mac: Applications > Utilities > Terminal
 - Windows: Download puTTY from <http://www.putty.org/>
 - Linux: You should know how to do this.

© 2017 Regents of the University of Minnesota. All rights reserved.



Part 1: Let's Access MSI

2. Connect to the login node

- Mac/Linux:

1. Type `ssh user@login.msi.umn.edu`
2. Type 'y' when the security alert shows up.
3. Enter password.

- Windows:

1. Type `login.msi.umn.edu` into the Host Name box, click “Open”
2. Enter username and password as with Mac

© 2017 Regents of the University of Minnesota. All rights reserved.

Part 1: Let's Access MSI

3. Request an interactive session

- You should see something like
`username@login03 [~] %`
before your cursor
- Type `isub`
- You may have to wait a minute or two.
Once prompt returns, you are running in
an interactive session

© 2017 Regents of the University of Minnesota. All rights reserved.

Part 1: Let's Access MSI

4. Load the R software module

- Type `module load R/3.3.3`
- Type `R`
- This should work with other modules:
 - Try `module load samtools/1.6`
 - Try `module load ncbi_blast+/2.2.29`
 - Try `module avail bwa` to view BWA versions

© 2017 Regents of the University of Minnesota. All rights reserved.



Part 1: MSI Organization

System status

<https://www.msi.umn.edu/systemstatus>

System usage

<https://s3.msi.umn.edu/pbsnodes/index.html>

Queues

<https://www.msi.umn.edu/queues>

Writing a job script

<https://www.msi.umn.edu/content/job-submission-and-scheduling-pbs-scripts>

© 2017 Regents of the University of Minnesota. All rights reserved.

Outline

- Part 1: Intro to MSI Systems
 - HPC systems, Storage, Queues, Modules
- Part 2: Parallelization and Task Arrays
 - GNU Parallel, PBS Task Arrays
- Part 3: Git and GitHub

© 2017 Regents of the University of Minnesota. All rights reserved.



Part 2: Parallelization

- Many analyses can be broken down into independent tasks
- Examples:
 - Mapping reads from 60 samples to a reference genome
 - Generating thousands of gene alignments for relative rates calculations

© 2017 Regents of the University of Minnesota. All rights reserved.



Part 2: Parallel Tools on MSI

- **GNU Parallel:** fancy Perl script that monitors processes and interleaves commands
- **TORQUE task arrays:** Generate an array of jobs, run ranges of them at once as separate PBS jobs

© 2017 Regents of the University of Minnesota. All rights reserved.



Part 2: Parallel Tools on MSI

- Both are useful and have pros/cons
- GNU Parallel:
 - Somewhat non-intuitive, but very powerful
 - Requires tricks for multi-node jobs
- Task arrays:
 - Easy to use if you understand bash arrays
 - Can use multiple nodes without tricks

© 2017 Regents of the University of Minnesota. All rights reserved.



Part 2: Which Parallel Tool to Use?

- It's really just personal preference.
- My strategy: use both.
 - Task arrays when there are \leq dozens of expensive/long tasks, where any of them might fail and need to be re-run.
 - Parallel when there are \geq hundreds small/cheap tasks

© 2017 Regents of the University of Minnesota. All rights reserved.



Part 2: Using Task Arrays

- Example: You want to align paired end reads from 60 samples to your reference genome. Each sample should be aligned with identical parameters.

© 2017 Regents of the University of Minnesota. All rights reserved.



Part 2: Using Task Arrays

1. Organize your data in a way that is easy to programmatically access.

— Example:

`/panfs/roc/scratch/konox006/reads/sample1_R1.fastq.gz`

`/panfs/roc/scratch/konox006/reads/sample1_R2.fastq.gz`

`/panfs/roc/scratch/konox006/reads/sample2_R1.fastq.gz`

`/panfs/roc/scratch/konox006/reads/sample2_R2.fastq.gz`

...

`/panfs/roc/scratch/konox006/reads/sample60_R1.fastq.gz`

`/panfs/roc/scratch/konox006/reads/sample61_R2.fastq.gz`

© 2017 Regents of the University of Minnesota. All rights reserved.

Part 2: Using Task Arrays

2. Inside your job script, generate bash arrays that hold the units of parallelization. In this example, it will be reads files

— Example:

```
SAMPLES_F=$(find /panfs/roc/scratch/konox006/reads  
-type f -name '*R1.fastq.gz' | sort -V))
```

```
SAMPLES_R=$(find /panfs/roc/scratch/konox006/reads  
-type f -name '*R2.fastq.gz' | sort -V))
```

© 2017 Regents of the University of Minnesota. All rights reserved.

Part 2: Using Task Arrays

3. Use the `${PBS_ARRAYID}` variable to access an individual element of the arrays. In this case, the forward and reverse reads of a sample.

- Example:

```
CURRENT_FWD=${SAMPLES_F[${PBS_ARRAYID}]}  
CURRENT_REV=${SAMPLES_R[${PBS_ARRAYID}]}  
SAMPLE=$(basename ${CURRENT_FWD} | cut -f 1 -d  
'_')
```

© 2017 Regents of the University of Minnesota. All rights reserved.

Part 2: Using Task Arrays

4. Build the command using the element that you fetched from the array.

— Example:

```
bwa mem \  
-t 8 -k 12 -M \  
/home/group/shared/references/ref.fa \  
{CURRENT_FWD} {CURRENT_REV} \  
> /panfs/roc/scratch/konox006/aln/{SAMPLE}.sam
```

© 2017 Regents of the University of Minnesota. All rights reserved.



Part 2: Using Task Arrays

5. Submit the array to the scheduler with `qsub` and the `-t` option. It can understand both comma separated values and ranges.

- Example:

```
qsub -t 0-59 bwa.sh
```

(if jobs 33 and 48 fail)

```
qsub -t 33,48 bwa.sh
```

© 2017 Regents of the University of Minnesota. All rights reserved.

Part 2: Task Array Notes

- You will get an email about each job starting and exiting (if you have configured the PBS script that way), so watch out.
- Example script here:
https://github.com/TomJKono/Presentations/blob/master/MSI-Git/Task_Array_Example.sh

© 2017 Regents of the University of Minnesota. All rights reserved.



Part 2: Task Array Notes

- `${PBS_ARRAYID}` is a TORQUE specific variable. It's automatically set by the scheduler when a job starts.
- If `qsub` is called without the `-t` option, the value of `${PBS_ARRAYID}` will be `0`
- Ranges are inclusive. `qsub -t 0-10` will start 11 jobs: 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10.

© 2017 Regents of the University of Minnesota. All rights reserved.

Part 2: Using GNU Parallel

- Example: You want to generate alignments for relative rates calculations. You have thousands of genes to align, but each alignment should take at most 15 minutes.

© 2017 Regents of the University of Minnesota. All rights reserved.



Part 2: Using GNU Parallel

1. Organize your data in a way that it is easy to programmatically access, similar to task arrays

- Example:

```
/panfs/roc/scratch/konox006/genes/gene1_group.fasta  
/panfs/roc/scratch/konox006/genes/gene2_group.fasta  
/panfs/roc/scratch/konox006/genes/gene3_group.fasta  
/panfs/roc/scratch/konox006/genes/gene4_group.fasta  
...  
/panfs/roc/scratch/konox006/genes/gene7999_group.fasta  
/panfs/roc/scratch/konox006/genes/gene8000_group.fasta
```

© 2017 Regents of the University of Minnesota. All rights reserved.



Part 2: Using GNU Parallel

2. Build the commands, and save them to a text file with `echo`

- Example:

```
for gene in /panfs/roc/scratch/konox006/genes/*  
do  
echo "clustalo -v -i ${gene} -t Protein \  
      --full-iter --iter=10 --out /scratch/${gene}"  
done > commands.txt
```

© 2017 Regents of the University of Minnesota. All rights reserved.

Part 2: Using GNU Parallel

3. Use `parallel` to dispatch these individual jobs to the multiple cores on the node. Some arithmetic with `ppn` may be necessary.

— Example:

```
#PBS -l mem=22gb,nodes=1:ppn=24,walltime=24:00:00
```

```
...
```

```
parallel --jobs 4 < commands.txt
```

© 2017 Regents of the University of Minnesota. All rights reserved.

Part 2: Using GNU Parallel

4. Some tricks are necessary for **multi-node** jobs

- See this page:
<https://www.msi.umn.edu/support/faq/how-can-i-use-gnu-parallel-run-lot-commands-parallel>
- Basically: modify the parallel environment, then use `--sshloginfile` and `${PBS_NODEFILE}` to dispatch jobs across nodes

© 2017 Regents of the University of Minnesota. All rights reserved.



Part 2: Let's Use Parallel

- Request an interactive session on MSI
`isub -m 4gb -w 2:00:00 -q lab`
- Load the parallel module
`module load parallel`

© 2017 Regents of the University of Minnesota. All rights reserved.



Part 2: Let's Use Parallel

- Make a new text file `commands.txt`
echo "a"
echo "b"
echo "c"
echo "d"
echo "this is a longer string"

© 2017 Regents of the University of Minnesota. All rights reserved.



Part 2: Let's Use Parallel

- Run it with bash, in serial
`bash commands.txt`

- Output:

`a`

`b`

`c`

`this is a longer string`

© 2017 Regents of the University of Minnesota. All rights reserved.

Part 2: Let's Use Parallel

- Now, run it with parallel
`parallel < commands.txt`

- Output (order may vary):

c

b

a

this is a longer string

© 2017 Regents of the University of Minnesota. All rights reserved.

Part 2: Parallel Tips

- Command lines are separated by newlines by default. This means you can string commands together to track progress for longer jobs
- You can even use the bash command separators `&&`, `| |`, and `;`

© 2017 Regents of the University of Minnesota. All rights reserved.



Part 2: Parallel Tips

- Example: Keep track of commands that succeed

```
clustalo ... && echo "gene1.fasta" >> Pass.txt
```

- Example: Keep track of commands that fail

```
clustalo ... || echo "gene1.fasta" >> Fail.txt
```

© 2017 Regents of the University of Minnesota. All rights reserved.

Part 2: Parallel Tips

- There is a LOT more you can do with `parallel`. See the manual page:
<https://www.gnu.org/software/parallel/man.html>
- Example `parallel` script here:
https://github.com/TomJKono/Presentations/blob/master/MSI-Git/Parallel_Example.sh

With commands here:

https://github.com/TomJKono/Presentations/blob/master/MSI-Git/Parallel_Commands.txt

© 2017 Regents of the University of Minnesota. All rights reserved.

Outline

- Part 1: Intro to MSI Systems
 - HPC systems, Storage, Queues, Modules
- Part 2: Parallelization and Task Arrays
 - GNU Parallel, PBS Task Arrays
- Part 3: Git and GitHub

© 2017 Regents of the University of Minnesota. All rights reserved.



Part 3: Why Git?

- If anything, to keep a working version of a script, in case you break it
- (Relatively) easy to collaborate with other scientists on the same script
- Great place to host scripts for reproducibility after publication

© 2017 Regents of the University of Minnesota. All rights reserved.



Part 3: Git and GitHub

- Git is the software that manages versions and allows users to “clone” and modify repositories
- GitHub is a service provider that hosts git servers and manages access to repositories

© 2017 Regents of the University of Minnesota. All rights reserved.



Part 3: UMN Has GitHub

- Specifically, we have Enterprise GitHub
 - Unlimited private repositories (can control who can see/contribute to a repo)
 - Only available to people with an X500 name and password
- <http://github.umn.edu/>

© 2017 Regents of the University of Minnesota. All rights reserved.



Part 3: Installing Git

- Download the software
- Mac:
<https://git-scm.com/download/mac>
- Windows:
<https://git-scm.com/download/win>
- Note: I don't know how to use git on windows, sorry

© 2017 Regents of the University of Minnesota. All rights reserved.



Part 3: Verify Installation

```
[22:09:00] [tomkono@Toms-MacBook-Pro] (~) $ which git
/usr/local/bin/git
[22:09:07] [tomkono@Toms-MacBook-Pro] (~) $ git
usage: git [--version] [--help] [-C <path>] [-c name=value]
         [--exec-path[=<path>]] [--html-path] [--man-path] [--info-path]
         [-p | --paginate | --no-pager] [--no-replace-objects] [--bare]
         [--git-dir=<path>] [--work-tree=<path>] [--namespace=<name>]
         <command> [<args>]
```

These are common Git commands used in various situations:

start a working area (see also: git help tutorial)

clone	Clone a repository into a new directory
init	Create an empty Git repository or reinitialize an existing one

work on the current change (see also: git help everyday)

add	Add file contents to the index
mv	Move or rename a file, a directory, or a symlink
reset	Reset current HEAD to the specified state
rm	Remove files from the working tree and from the index

© 2017 Regents of the University of Minnesota. All rights reserved.

Part 3: Let's Use UMN GitHub

- Sign in to <http://github.umn.edu/> to activate your account

GitHub Enterprise

Sign in via LDAP

Username

konox006

Password

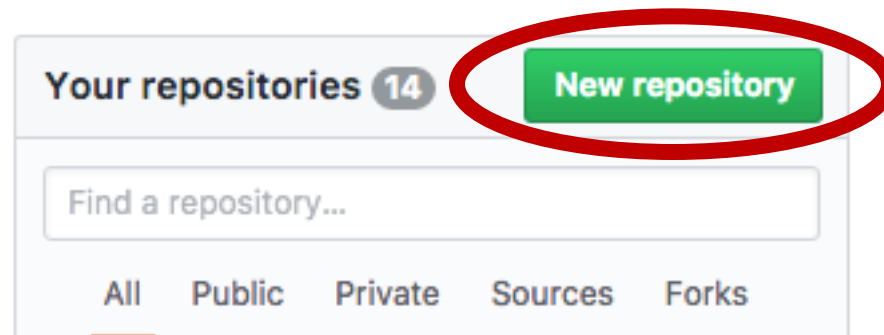
.....

Sign in

© 2017 Regents of the University of Minnesota. All rights reserved.

Part 3: Let's Use UMN GitHub

- Start a new repository with this button



© 2017 Regents of the University of Minnesota. All rights reserved.


Part 3: Let's Use UMN GitHub

- Fill in the details

Create a new repository

A repository contains all the files for your project, including the revision history.

Owner

 konox006 ▾

Repository name

Nature_in_prep

Great repository names are short and memorable. Need inspiration? How about **silver-funicular**.

Description (optional)

My hella cool pub, in prep



Public

Any logged in user can see this repository. You choose who can commit.



Private

You choose who can see and commit to this repository.

☒ **Initialize this repository with a README**

This will let you immediately clone the repository to your computer. Skip this step if you're importing an existing repository.

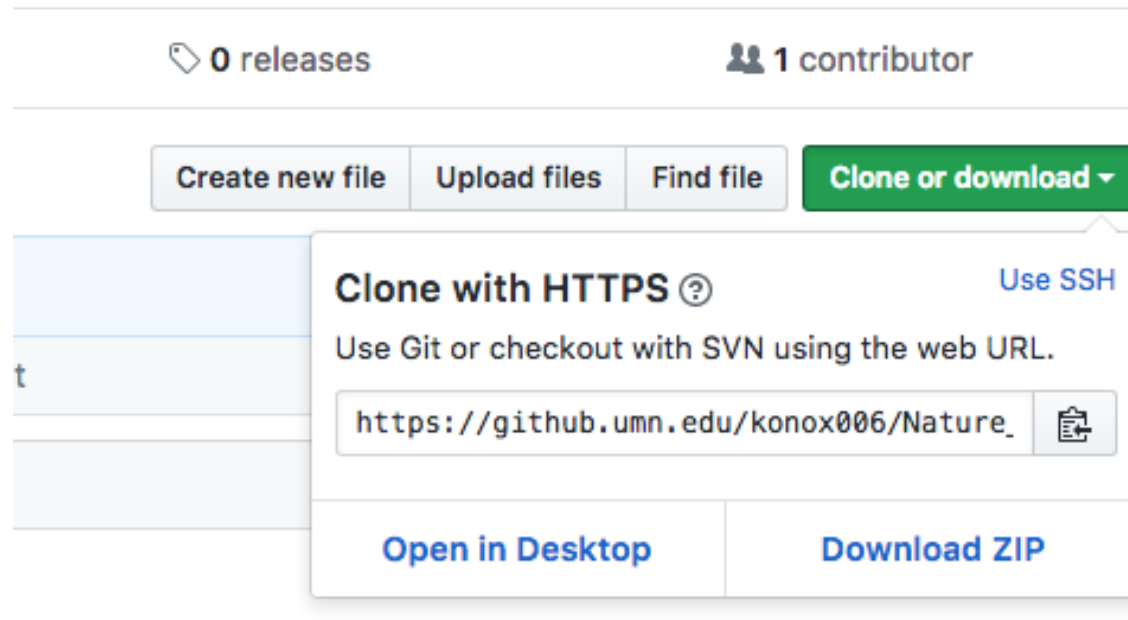
Add .gitignore: **None** ▾

Create repository

© 2017 Regents of the University of Minnesota. All rights reserved.

Part 3: Let's Use UMN GitHub

- Show the “clone” button to get the clone URL



© 2017 Regents of the University of Minnesota. All rights reserved.

Part 3: Let's Use UMN GitHub

- Clone the repo in your Terminal window

```
[22:35:51] [tomkono@Toms-MacBook-Pro] (~) $ git clone https://github.umn.edu/konox006/Nature_in_prep.git
Cloning into 'Nature_in_prep'...
remote: Counting objects: 3, done.
remote: Total 3 (delta 0), reused 0 (delta 0), pack-reused 0
Unpacking objects: 100% (3/3), done.
[22:36:05] [tomkono@Toms-MacBook-Pro] (~) $ cd Nature_in_prep/
[22:36:15] [tomkono@Toms-MacBook-Pro] (~/.Nature_in_prep) $ █
```

© 2017 Regents of the University of Minnesota. All rights reserved.



Part 3: Let's Use UMN GitHub

- Edit the README.md file to add some additional detail

```
1 # Nature in prep
2 My hella cool pub, in prep
3
4 This research project is so great, it'll knock your socks off. Trust me, even
5 though I haven't done any actual analyses or collected any data, this stuff
6 will be tip-top.
```

© 2017 Regents of the University of Minnesota. All rights reserved.

Part 3: Let's Use UMN GitHub

- View the changes with `status`, and `add` them to the repository

```
[22:48:59] [tomkono@Toms-MacBook-Pro] (~/.Nature_in_prep) $ git status
On branch master
Your branch is up to date with 'origin/master'.

Changes not staged for commit:
  (use "git add <file>..." to update what will be committed)
  (use "git checkout -- <file>..." to discard changes in working directory)

        modified:   README.md

no changes added to commit (use "git add" and/or "git commit -a")
[22:49:01] [tomkono@Toms-MacBook-Pro] (~/.Nature_in_prep) $ git add README.md
```

© 2017 Regents of the University of Minnesota. All rights reserved.

Part 3: Let's Use UMN GitHub

- Then `commit` them to create a “save point” in the repo history. Use `-m` to add a “commit message” that will display in the repository history. Make it meaningful (more than just “added stuff”)

```
[22:49:04] [tomkono@Toms-MacBook-Pro] (~/Nature_in_prep) $ git commit -m 'Put crap in the README'  
[master 8e51dab] Put crap in the README  
1 file changed, 4 insertions(+)
```

© 2017 Regents of the University of Minnesota. All rights reserved.



Part 3: Let's Use UMN GitHub

- Actually, that commit message is not very nice. Let's **amend** it:

```
[22:58:27] [tomkono@Toms-MacBook-Pro] (~/.Nature_in_prep) $ git commit --amend
```

```
1 Put additional text in the README
2
3 # Please enter the commit message for your changes. Lines starting
4 # with '#' will be ignored, and an empty message aborts the commit.
5 #
6 # Date:      Sun Nov 12 22:58:27 2017 -0600
7 #
8 # On branch master
9 # Your branch is ahead of 'origin/master' by 1 commit.
10 # (use "git push" to publish your local commits)
11 #
12 # Changes to be committed:
13 #>   modified:   README.md
14 #
```

© 2017 Regents of the University of Minnesota. All rights reserved.

Part 3: Let's Use UMN GitHub

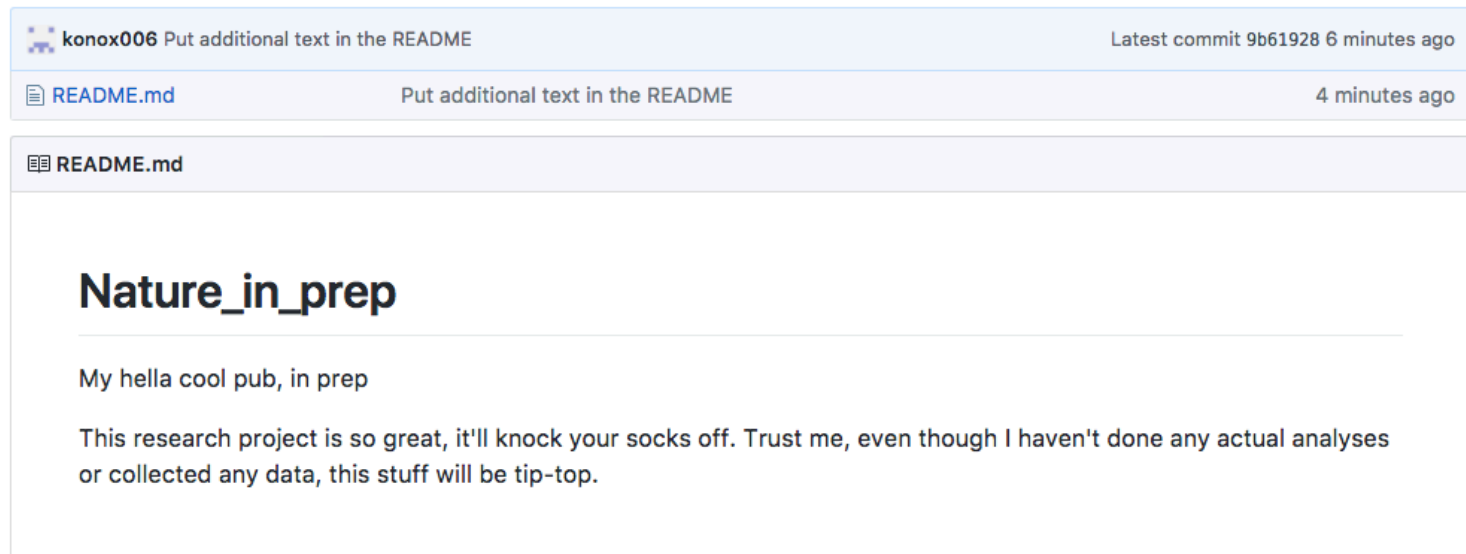
- Now, we are ready to **push** the changes to the repository

```
[23:02:04] [tomkono@Toms-MacBook-Pro] (~/Nature_in_prep) $ git push
Counting objects: 3, done.
Delta compression using up to 4 threads.
Compressing objects: 100% (2/2), done.
Writing objects: 100% (3/3), 422 bytes | 422.00 KiB/s, done.
Total 3 (delta 0), reused 0 (delta 0)
To https://github.umn.edu/konox006/Nature_in_prep.git
   2a8a8b0..9b61928  master -> master
[23:02:07] [tomkono@Toms-MacBook-Pro] (~/Nature_in_prep) $
```

© 2017 Regents of the University of Minnesota. All rights reserved.

Part 3: Let's Use UMN GitHub

- The new commit will appear on the repo website



© 2017 Regents of the University of Minnesota. All rights reserved.

Part 3: Let's Use UMN GitHub

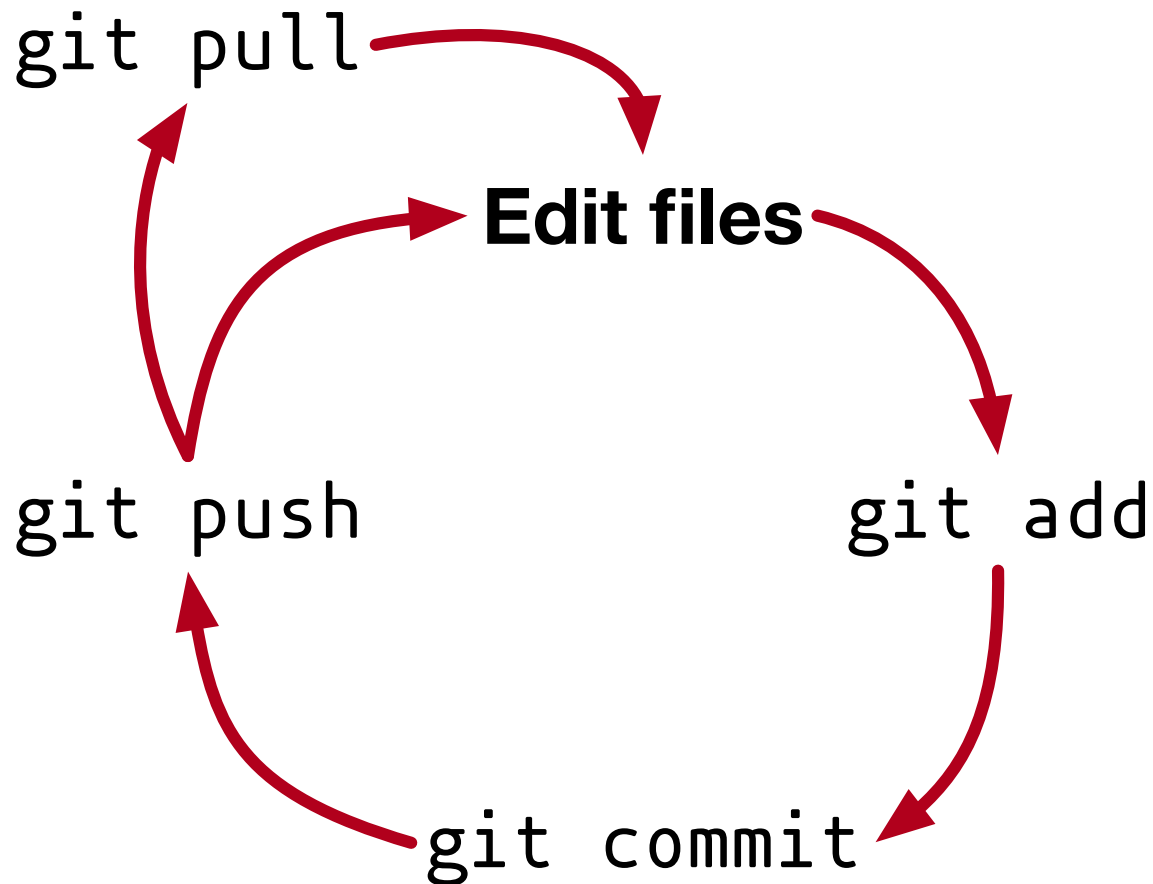
- When working with others, be sure to **pull** remote changes before editing files, else you'll create a "merge conflict" which must be **manually resolved**.
- **This is painful. Avoid it.**
- **I had one with LaTeX and I just burned the repo and re-cloned it**

© 2017 Regents of the University of Minnesota. All rights reserved.



Part 3: Let's Use UMN GitHub

- Put your own script into the repo, and go through `add/commit/push`



© 2017 Regents of the University of Minnesota. All rights reserved.

Part 3: Git Resources

- GitHub cheat sheet
<https://education.github.com/git-cheat-sheet-education.pdf>
- Interactive git tutorial
<https://try.github.io/>
- Vince Buffalo's book (Chapter 5)
<http://vincebuffalo.org/book/>

© 2017 Regents of the University of Minnesota. All rights reserved.



Thanks! Questions?

- What can I do better?

The University of Minnesota is an equal opportunity educator and employer. This presentation is available in alternative formats upon request. Direct requests to Minnesota Supercomputing Institute, 599 Walter Library, 117 Pleasant St. SE, Minneapolis, Minnesota, 55455, 612-624-0528.

© 2017 Regents of the University of Minnesota. All rights reserved.

